

Hyperbolic Tangent $g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$
 $g'(z) = 1 - g(z)^2$
 Rectified Linear Unit (ReLU)
 $g(z) = \max(0, z)$ $g'(z) = \begin{cases} 1, & z > 0 \\ 0, & \text{otherwise} \end{cases}$

- Local connectivity of hidden units
- share weights across certain units
 - pool hidden units in the same neighborhood
 - interleave feature extraction and pooling

Gradient Descent

- Pure analytically: give symbolic answer, hard to compute.
 - Finite difference approximation: $O(NM)$
 - Back Propagation: exact answer, hard to implement $O(NM)$
- Find $g = b + c$, $b = aw_2$, $a = w_1$, $c = 3w_1$, 这样算.

Regression

Linear Regression

prediction $hw(x) = w_0 + w_1 x$

Loss function: $L(w) = \sum_i (y_i - hw(x_i))^2 = \sum_i (y_i - w^T x_i)^2$

$w^* = (X^T X)^{-1} X^T y$

Regularization alleviate overfitting

LASSO (Least absolute shrinkage and selection)

$L(w) = \sum_i (y_i - w^T x_i)^2 + \lambda \sum_k |w_k|$

Ridge Regression (L_2)

$L(w) = \sum_i (y_i - w^T x_i)^2 + \lambda \sum_k w_k^2$

Unsupervised ML

1. k-means

the point y with min squared Euclidean distance to a set of point is their mean.

问题: overlap clusters, 有些 cluster 会重叠, data 会

2. Expectation-Maximization (EM)

① Gaussian Mixture Model (GMM) $\pi \rightarrow \mu, \sigma^2$

$p(x) = \sum_{k=1}^K \pi_k N(x | \mu_k, \Sigma_k)$

$\forall k: \pi_k > 0, \sum_{k=1}^K \pi_k = 1$

$P(Y)$: Distribution over k components (clusters)

$P(X|Y)$: Each component generates data from a multi-var

Gaussian with μ_k and Σ_k .

Each point is sampled from a generative process.

① choose component i with $p = \pi_i$ (random)

② Generate data-point from $N(x | \mu_i, \Sigma_i)$

③ Don't know label y .

$\max_{\theta} L = \prod_{j=1}^n P(x_j) = \prod_{j=1}^n \sum_{i=1}^K \pi_i P(y_j = i, x_j) = \prod_{j=1}^n \sum_{i=1}^K \pi_i N(x_j | \mu_i, \Sigma_i)$

④ EM (Expectation-Maximization)

• Pick k random cluster models (Gaussians)

• Loop until convergence:

• Assign data instances proportionately to different

models

• Revise each model based on its points

Object Function: $\arg \max_{\theta} \prod_{j=1}^n \sum_{i=1}^K \pi_i P(y_j = i, x_j | \theta)$

Data: $\{x_j | j = 1 \dots n\}$

E-step: compute Expectation to fill in missing y values

according to current para: θ .

\Rightarrow for sample j , $P(y_j = i | x_j, \theta)$.

M-step: reestimate the para with weighted MLE

$\text{set } \theta = \arg \max_{\theta} L(\theta)$

⑤ EM \rightarrow k-means when

① All Gaussians are spherical, w_i, Σ_i same. ($\theta = \{\mu_i\}$)

② hard assignment of points to Gaussians. ($\Sigma_i \rightarrow 0$)

$Q(s_t, a_t) \leftarrow (1-\alpha) Q(s_t, a_t) + \alpha (r_t + \gamma \max_{a'} Q(s_{t+1}, a'))$

Markov Decision Process

1. MDP Model
 a set of states S , a set of actions A ,
 a transition function $T(s, a, s') = P(s' | s, a)$,
 a reward function $R(s, a, s') = R(s, s')$,
 a start state, sometimes a terminal state.

2. Utilities

Policy $\pi: S \rightarrow A$ π^* optimize expected utilities.

Utility of states: $V^*(s)$ = expected utility starting in s and acting optimally

Utility of Q state (s, a) : $Q^*(s, a)$: $s \rightarrow \{s'; \text{if act optimal to exp utility}\}$

$V^*(s) = \max_a Q^*(s, a)$

$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$

$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$

3. Value Iteration $\pi^*(s, a)$ Policy π value 变 π converge (max only change)

$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$

$\forall s$ converge to optimal (unique).

4. Policy extraction

$\pi^*(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$

$\pi^*(s) = \arg \max_a Q^*(s, a)$ (more useful, compute fast)

5. Policy Iteration optimal: converge faster, repeat ②④

① Policy evaluation $Q(s, a)$ 每次 iteration

a. iterative update

$V_0(s) = 0, V_{k+1}(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k(s')]$

直到 converge. 每次迭代 $Q(s, \pi(s))$ [converge, 但不 optimal]

b. π max. Bellman equation 可直接用 linear system solver 解.

$V^*(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^*(s')]$

② Policy improvement one step look ahead:

$\pi_{k+1}(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$

6. 问题

Reinforcement Learning

1. Model: MDP π, T, R, γ on \mathcal{H} .

2. Model-based learning

① Learn Empirical MDP: \hat{T}, \hat{R} 估计 T, R

② $\hat{\pi}$ MDP.

3. Model Free Learning

① Passive RL: policy evaluation. values of

• Direct evaluation: goal: compare states under π .

Idea: average over samples.

• π policy π 是, everytime you visit a state, write

down what the sum of discounted rewards turned out

Pro: eventually compute correct average values.

Con: waste info about state connections, each state should

be learned separately, time long. experience

• Temporal Difference Learning: learn Q from every

sample $= R(s, \pi(s), s') + \gamma V^*(s')$

$V^*(s) \leftarrow (1-\alpha) V^*(s) + \alpha \text{sample} = V^*(s) + \alpha (\text{sample} - V^*(s))$

• Exponential moving avg

$\bar{x}_n = \frac{x_n + (1-\alpha)x_{n-1} + (1-\alpha)^2 x_{n-2} + \dots}{1 + (1-\alpha) + (1-\alpha)^2 + \dots}$ $\alpha \downarrow$, converge.

② Passive RL: Q -learning

Learn $Q(s, a)$ as you go:

• Receive a sample (s, a, s', r)

• consider old estimate $Q(s, a)$

• consider new estimate: sample $= R(s, a, s') + \gamma \max_{a'} Q(s', a')$

$Q(s, a) \leftarrow (1-\alpha) Q(s, a) + \alpha (\text{sample})$

off-policy planning: converges to optimal policy even

acting suboptimally. [Explore enough, 别太 slow]

4. Exploration v.s. Exploitation

① ϵ -Greedy: random actions with ϵ , policy act $1-\epsilon$.

问题: explore the space but thrashing around.

solution: ϵ reduction / exploration functions

② Exploration Functions: select actions based on R and Q value

$f(a, n) = u + \frac{k}{n}$ u : Q -value estimation, n : visit count, returns an optimistic utility

$Q(s, a) \leftarrow \alpha R(s, a, s') + \gamma \max_{a'} f(a(s', a'), N(s', a'))$

5. Regret: 最优的 Q 和 R 政策, $R-R^*$

ϵ greedy $>$ exploration-function.

Approximate Q-learning

$V(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$

$Q(s, a) = w_1 f_1(s, a) + \dots + w_n f_n(s, a)$

$Q(s, a) \leftarrow Q(s, a) + \alpha [\text{difference}]$

$w_i \leftarrow w_i + \alpha [\text{difference}] f_i(s, a)$

difference $= [r + \gamma \max_{a'} Q(s', a')] - Q(s, a)$

transition $= (s, a, r, s')$

Policy Search

• Observation: often the feature-based

policies that work well (win games,

maximize utilities) aren't the ones that

approximate V / Q best

• Idea: learn policies that maximize rewards,

not the values that predict them

• Policy search: start with an OK solution (e.g.,

approximate Q -learning), then fine-tune

feature weights to find a better policy

• Simplest policy search:

• Start with an initial linear value function or

Q -function

• Change each feature weight up and down

and see if your policy is better than before

Problems:

• How do we tell the policy got better?

• Need to run many sample episodes!

• If there are a lot of features, this can be

impractical

Dependency Parse

Advantages

• Deals well with free word order languages

where the constituent structure is quite fluid

• Ex: Czech, Turkish

• Parsing is much faster than CFG-bases

parsers

• Dependency structure often captures the

syntactic relations needed by later

applications

• CFG-based approaches often extract

this same information from trees anyway

Dependency Parsing

• Parsing: taking a string and a grammar and

returning one or more parse tree(s) for that

string

• There are two modern approaches to

dependency parsing

• Graph-based approach: finding the

(maximum) spanning trees of the complete

graph over words

• Transition-based (shift-reduce)

approach: reading words from left to right

and taking a sequence of actions to

construct a tree

EM in General

• Can be used to learn any model with hidden

variables (missing data)

• Alternate:

• Compute distributions over hidden

variables based on current parameter values

• Compute new parameter values to

maximize expected log likelihood based on

distributions over hidden variables

• Stop when no changes

• EM degrades to k-means if we assume

• All the Gaussians are spherical and have

identical weights and covariances

• i.e., the only parameters are the means

• The label distributions computed at E-

step are point-estimations

• i.e., hard-assignments of data points to

Gaussians

• Alternatively, assume the variances are

close to zero

Transition matrix

X_{t+1} $P(X_{t+1} | X_t)$

$\begin{matrix} & \text{sun} & \text{rain} \\ \text{sun} & 0.9 & 0.1 \\ \text{rain} & 0.3 & 0.7 \end{matrix} \rightarrow T = \begin{pmatrix} 0.9 & 0.3 \\ 0.1 & 0.7 \end{pmatrix}$

$P(X_t) = 0.6 < 0.9, 0.1 > + 0.4 < 0.3, 0.7 >$

$= T \begin{pmatrix} 0.6 \\ 0.4 \end{pmatrix}$